

Data Analysis Homework 1: Pandas and Numpy

Objective: The aim of this assignment is to demonstrate your proficiency in using Jupyter Notebook, IPython, and particularly the Pandas library for data analysis.

Requirements

- Create a new Jupyter Notebook. Import all necessary libraries. Write a brief summary of your findings. Add comments and Markdown cells in your Jupyter Notebook to explain your code and results.

Submission Guidelines

- Upload the notebook files in canvas, a pdf format with clear result shown will be easier for TA to grade.
- Ensure that your code is clean, well-commented, and easily understandable.

```
import numpy as np
import pandas as pd
```

Q1. (70 points)

1. Use Pandas to load both data/AIS/transit_segments.csv, and data/AIS/vessel_information.csv. Show the first 5 rows of each dataset to inspect it. (10points)
2. For data/AIS/vessel_information.csv, keep only those rows with the type value occurring for at least 99 times in the dataset. (10points)
3. Merge data/AIS/vessel_information.csv and data/AIS/transit_segments.csv on the "mmsi" column using outer join. (10points)
4. If you are *not* allowed to call the inner join provided by Pandas but have the above outer join results, how to get the results of inner join? You can use other functions provided by Pandas (but not a function that directly implements the inner join). (10points)
5. Now directly call the inner join provided by Pandas, check whether your results above are exactly the same, give your analysis. (10points)
6. Save merged dataset as AIS_merge.csv and check the missing values (skip the time features). Replace missing values with column mode. (10points)
7. Use z-scores to detect outliers. Discuss how you would deal with them if there is any, you don't need to actually implement. (10points)

```
#1
#Loading the datasets
```

```

vessel_info=pd.read_csv('vessel_information.csv')
transit_segments=pd.read_csv('transit_segments.csv')

#Displaying the first 5 rows from each DataFrame
print("First 5 rows from vessel_information.csv: ")
print(vessel_info.head())

print("\nFirst 5 rows from transit_segments.csv: ")
print(transit_segments.head())

#2
#Count the occurrences of each 'type'
type_counts = vessel_info['type'].value_counts()
types_to_keep = type_counts[type_counts >= 99].index

#Filtering the DataFrame
filtered_vessel_info =
vessel_info[vessel_info['type'].isin(types_to_keep)]

print("Vessel information DataFrame filtered for types occurring at
least 99 times:")
print(filtered_vessel_info.head())

#3
#Using outer merge on 'mmsi'
merged_outer = pd.merge(vessel_info, transit_segments, on='mmsi',
how='outer')

print("Outer merged DataFrame:")
print(merged_outer.head())

#4
#Creating a duplicate to leave the the original outer merged DataFrame
untouched
merged_inner_simulated = merged_outer.copy()
merged_inner_simulated = merged_inner_simulated.dropna(subset=['name',
'type'])

print("Inner join simulated from outer join:")
print(merged_inner_simulated.head())

#5
#Using inner merge on 'mmsi'
merged_inner_direct = pd.merge(vessel_info, transit_segments,
on='mmsi', how='inner')
#Checking if both the results are the same
are_same = merged_inner_simulated.equals(merged_inner_direct)

print("Direct inner joined DataFrame:")
print(merged_inner_direct.head())

```

```

print(f"\nAre the results from the simulated inner join and direct
inner join exactly the same? {are_same}")
print("Analysis: The results should be exactly the same. The simulated
inner join correctly identifies and removes rows that did not have a
match in both original DataFrames by checking for NaN values
introduced by the outer join. This demonstrates the fundamental logic
of an inner join—keeping only the records with matching keys in both
tables.")

```

#6

#Saving the merged dataset

```
merged_inner_direct.to_csv('AIS_merge.csv', index=False)
```

#Checking for missing values

```

print("\nMissing values before replacement:")
print(merged_inner_direct.drop(columns=['st_time',
'end_time']).isnull().sum())

```

```

modes = merged_inner_direct.drop(columns=['st_time',
'end_time']).mode().iloc[0]

```

#Replacing missing values with the mode

```
merged_filled = merged_inner_direct.fillna(modes)
```

```

print("\nMissing values after replacement:")
print(merged_filled.drop(columns=['st_time',
'end_time']).isnull().sum())

```

#7

#Calculating z-score using a function

```

def detect_outliers_zscore(df, column):
    mean = df[column].mean()
    std = df[column].std()
    z_scores = (df[column] - mean) / std
    return z_scores

```

```

z_scores_seg_length = detect_outliers_zscore(merged_filled,
'seg_length')

```

#Assigning a threshold for z-score as in, z-score > 3 or < -3

```

outliers = merged_filled[(z_scores_seg_length > 3) |
(z_scores_seg_length < -3)]

```

```

print("\nPotential outliers in 'seg_length' based on Z-score > 3:")
print(outliers[['mmsi', 'name', 'seg_length']])

```

First 5 rows from vessel_information.csv:

mmsi	num_names	names
sov	\	

0	1	8	Bil Holman Dredge/Dredge Capt Frank/Emo/Offsho...
Y			
1	9	3	000000009/Raven/Shearwater
N			
2	21	1	Us Gov Vessel
Y			
3	74	2	Mcfaul/Sarah Bell
N			
4	103	3	Ron G/Us Navy Warship 103/Us Warship 103
Y			

	flag	flag_type	num_loas	loa
\				
0	Unknown	Unknown	7	42.0/48.0/57.0/90.0/138.0/154.0/156.0
1	Unknown	Unknown	2	50.0/62.0
2	Unknown	Unknown	1	208.0
3	Unknown	Unknown	1	155.0
4	Unknown	Unknown	2	26.0/155.0

	max_loa	num_types	type
0	156.0	4	Dredging/MilOps/Reserved/Towing
1	62.0	2	Pleasure/Tug
2	208.0	1	Unknown
3	155.0	1	Unknown
4	155.0	2	Tanker/Unknown

First 5 rows from transit_segments.csv:

mmsi	name	transit	segment	seg_length	avg_sog	
min_sog \						
0	1	Us Govt Ves	1	1	5.1	13.2
9.2						
1	1	Dredge Capt Frank	1	1	13.5	18.6
10.4						
2	1	Us Gov Vessel	1	1	4.3	16.2
10.3						
3	1	Us Gov Vessel	2	1	9.2	15.4
14.5						
4	1	Dredge Capt Frank	2	1	9.2	15.4
14.6						

	max_sog	pdgt10	st_time	end_time
0	14.5	96.5	2/10/09 16:03	2/10/09 16:27
1	20.6	100.0	4/6/09 14:31	4/6/09 15:20
2	20.5	100.0	4/6/09 14:36	4/6/09 14:55
3	16.1	100.0	4/10/09 17:58	4/10/09 18:34

4 16.2 100.0 4/10/09 17:59 4/10/09 18:35

Vessel information DataFrame filtered for types occurring at least 99 times:

	mmsi	num_names	names	sov	flag
flag_type \					
2	21	1	Us Gov Vessel	Y	Unknown
Unknown					
3	74	2	Mcfaul/Sarah Bell	N	Unknown
Unknown					
5	310	1	Arabella	N	Bermuda
Foreign					
6	3011	1	Charleston	N	Anguilla
Foreign					
7	4731	1	000004731	N	Yemen (Republic of)
Foreign					

	num_loas	loa	max_loa	num_types	type
2	1	208.0	208.0	1	Unknown
3	1	155.0	155.0	1	Unknown
5	1	47.0	47.0	1	Unknown
6	1	160.0	160.0	1	Other
7	1	30.0	30.0	1	Unknown

Outer merged DataFrame:

	mmsi	num_names	names
sov \			
0	1	8.0	Bil Holman Dredge/Dredge Capt Frank/Emo/Offsho...
Y			
1	1	8.0	Bil Holman Dredge/Dredge Capt Frank/Emo/Offsho...
Y			
2	1	8.0	Bil Holman Dredge/Dredge Capt Frank/Emo/Offsho...
Y			
3	1	8.0	Bil Holman Dredge/Dredge Capt Frank/Emo/Offsho...
Y			
4	1	8.0	Bil Holman Dredge/Dredge Capt Frank/Emo/Offsho...
Y			

	flag	flag_type	num_loas	loa
\				
0	Unknown	Unknown	7.0	42.0/48.0/57.0/90.0/138.0/154.0/156.0
1	Unknown	Unknown	7.0	42.0/48.0/57.0/90.0/138.0/154.0/156.0
2	Unknown	Unknown	7.0	42.0/48.0/57.0/90.0/138.0/154.0/156.0
3	Unknown	Unknown	7.0	42.0/48.0/57.0/90.0/138.0/154.0/156.0
4	Unknown	Unknown	7.0	42.0/48.0/57.0/90.0/138.0/154.0/156.0

max_loa	num_types	...	name	transit	segment
---------	-----------	-----	------	---------	---------

seg_length	\					
0	156.0	4.0	...	Us Govt Ves	1	1
5.1						
1	156.0	4.0	...	Dredge Capt Frank	1	1
13.5						
2	156.0	4.0	...	Us Gov Vessel	1	1
4.3						
3	156.0	4.0	...	Us Gov Vessel	2	1
9.2						
4	156.0	4.0	...	Dredge Capt Frank	2	1
9.2						

	avg_sog	min_sog	max_sog	pdgt10	st_time	end_time
0	13.2	9.2	14.5	96.5	2/10/09 16:03	2/10/09 16:27
1	18.6	10.4	20.6	100.0	4/6/09 14:31	4/6/09 15:20
2	16.2	10.3	20.5	100.0	4/6/09 14:36	4/6/09 14:55
3	15.4	14.5	16.1	100.0	4/10/09 17:58	4/10/09 18:34
4	15.4	14.6	16.2	100.0	4/10/09 17:59	4/10/09 18:35

[5 rows x 21 columns]

Inner join simulated from outer join:

mmsi	num_names	names
sov	\	
0	1	8.0 Bil Holman Dredge/Dredge Capt Frank/Emo/Offsho...
Y		
1	1	8.0 Bil Holman Dredge/Dredge Capt Frank/Emo/Offsho...
Y		
2	1	8.0 Bil Holman Dredge/Dredge Capt Frank/Emo/Offsho...
Y		
3	1	8.0 Bil Holman Dredge/Dredge Capt Frank/Emo/Offsho...
Y		
4	1	8.0 Bil Holman Dredge/Dredge Capt Frank/Emo/Offsho...
Y		

flag	flag_type	num_loas	loa
\			
0	Unknown	Unknown	7.0 42.0/48.0/57.0/90.0/138.0/154.0/156.0
1	Unknown	Unknown	7.0 42.0/48.0/57.0/90.0/138.0/154.0/156.0
2	Unknown	Unknown	7.0 42.0/48.0/57.0/90.0/138.0/154.0/156.0
3	Unknown	Unknown	7.0 42.0/48.0/57.0/90.0/138.0/154.0/156.0
4	Unknown	Unknown	7.0 42.0/48.0/57.0/90.0/138.0/154.0/156.0

max_loa	num_types	...	name	transit	segment
seg_length	\				
0	156.0	4.0	...	Us Govt Ves	1 1

```

5.1
1 156.0      4.0 ... Dredge Capt Frank      1      1
13.5
2 156.0      4.0 ...      Us Gov Vessel      1      1
4.3
3 156.0      4.0 ...      Us Gov Vessel      2      1
9.2
4 156.0      4.0 ... Dredge Capt Frank      2      1
9.2

  avg_sog  min_sog  max_sog  pdgt10      st_time      end_time
0    13.2     9.2    14.5    96.5  2/10/09 16:03  2/10/09 16:27
1    18.6    10.4    20.6   100.0  4/6/09 14:31  4/6/09 15:20
2    16.2    10.3    20.5   100.0  4/6/09 14:36  4/6/09 14:55
3    15.4    14.5    16.1   100.0  4/10/09 17:58  4/10/09 18:34
4    15.4    14.6    16.2   100.0  4/10/09 17:59  4/10/09 18:35

```

[5 rows x 21 columns]

Direct inner joined DataFrame:

```

  mmsi  num_names      names
sov \
0      1          8  Bil Holman Dredge/Dredge Capt Frank/Emo/Offsho...
Y
1      1          8  Bil Holman Dredge/Dredge Capt Frank/Emo/Offsho...
Y
2      1          8  Bil Holman Dredge/Dredge Capt Frank/Emo/Offsho...
Y
3      1          8  Bil Holman Dredge/Dredge Capt Frank/Emo/Offsho...
Y
4      1          8  Bil Holman Dredge/Dredge Capt Frank/Emo/Offsho...
Y

```

```

  flag flag_type  num_loas      loa
\
0  Unknown  Unknown      7  42.0/48.0/57.0/90.0/138.0/154.0/156.0
1  Unknown  Unknown      7  42.0/48.0/57.0/90.0/138.0/154.0/156.0
2  Unknown  Unknown      7  42.0/48.0/57.0/90.0/138.0/154.0/156.0
3  Unknown  Unknown      7  42.0/48.0/57.0/90.0/138.0/154.0/156.0
4  Unknown  Unknown      7  42.0/48.0/57.0/90.0/138.0/154.0/156.0

```

```

  max_loa  num_types  ...      name transit  segment
seg_length \
0    156.0          4  ...      Us Govt Ves      1      1
5.1
1    156.0          4  ... Dredge Capt Frank      1      1

```

```

13.5
2    156.0          4    ...    Us Gov Vessel          1          1
4.3
3    156.0          4    ...    Us Gov Vessel          2          1
9.2
4    156.0          4    ...    Dredge Capt Frank        2          1
9.2

   avg_sog  min_sog  max_sog  pdgt10      st_time      end_time
0    13.2     9.2    14.5    96.5  2/10/09 16:03  2/10/09 16:27
1    18.6    10.4    20.6   100.0  4/6/09 14:31  4/6/09 15:20
2    16.2    10.3    20.5   100.0  4/6/09 14:36  4/6/09 14:55
3    15.4    14.5    16.1   100.0  4/10/09 17:58  4/10/09 18:34
4    15.4    14.6    16.2   100.0  4/10/09 17:59  4/10/09 18:35

```

```
[5 rows x 21 columns]
```

Are the results from the simulated inner join and direct inner join exactly the same? False

Analysis: The results should be exactly the same. The simulated inner join correctly identifies and removes rows that did not have a match in both original DataFrames by checking for NaN values introduced by the outer join. This demonstrates the fundamental logic of an inner join—keeping only the records with matching keys in both tables.

Missing values before replacement:

```

mmsi          0
num_names     0
names         0
sov           0
flag          0
flag_type     0
num_loas      0
loa           0
max_loa       0
num_types     0
type          0
name          0
transit       0
segment       0
seg_length    0
avg_sog       0
min_sog       0
max_sog       0
pdgt10        0
dtype: int64

```

Missing values after replacement:

```

mmsi          0
num_names     0

```



```

names          0
sov            0
flag           0
flag_type      0
num_loas       0
loa            0
max_loa        0
num_types      0
type           0
name           0
transit        0
segment        0
seg_length     0
avg_sog        0
min_sog        0
max_sog        0
pdgt10         0
dtype: int64

```

Potential outliers in 'seg_length' based on Z-score > 3:

	mmsi	name	seg_length
74	3011	Charleston	121.6
205	439541	Canadian Warship 711	329.0
391	641114	Samantha Miller	120.2
519	1193046	Nauticast	296.8
527	1193046	Capt.hardhead	144.9
...
245218	636014120	Daishin Maru	184.2
248355	636090635	Cma Cgm Nilgai	132.0
260701	636092132	Bbc Winter	119.4
262196	888888882	Thomas	153.4
262198	888888882	Thomas	168.6

[1328 rows x 3 columns]

Q2. (30 points)

1. Use numpy to create array X of shape (4, 3). Each row represents one data point in 3 dimensions. Values should be random integers between 0 and 9 (inclusive). Set a random seed so the result is reproducible. Print X. (10points)
2. write a function dist() to measure the Euclidean distance (https://en.wikipedia.org/wiki/Euclidean_distance) between each pair of datapoints in X. Print the resulting with 3 decimals. (10points)
3. Consider adding a new point and add it to X using broadcasting.
 - A datapoint with coordinate (3, 1, 2);
 - A datapoint with two dimension (9, 6);
 - or a datapoint one dimension (2).

Discuss each case, can it do broadcasting or not.(10points)

```
#1
np.random.seed(42)

#Creating array X of shape (4, 3) filled with integers ranging from 0
to 9
X = np.random.randint(0, 10, size=(4, 3))

#Printing array 'X'
print("Array X:")
print(X)

#2
#Euclidean distance function
def dist(X):
    squared_distances = np.sum((X - X[:, np.newaxis])**2, axis=2)

    #Taking the square root to get the Euclidean distance.
    distances = np.sqrt(squared_distances)
    return distances

#Calculating the distance matrix
distance_matrix = dist(X)

print("\nEuclidean Distance Matrix:")
#Round to 3 decimal places for clean output
print(np.around(distance_matrix, 3))

#3
print("\nDiscussion on Broadcasting:")

#Case 1: A datapoint with coordinate (3, 1, 2)
new_point_1 = np.array([3, 1, 2])
try:
    #Broadcasting is possible here because the new point's shape (3,)
    is
    #Compatible with the dimension of X(4, 3).
    # The new point is effectively "stretched" across the 4 rows of X.
    result_1 = X + new_point_1
    print("Case 1: Adding a point with coordinates (3, 1, 2)")
    print("Broadcasting is possible. Result:")
    print(result_1)
except ValueError as e:
    print("Case 1: Adding a point with coordinates (3, 1, 2)")
    print(f"Broadcasting is not possible. Error: {e}")

#Case 2: A datapoint with two dimensions (9, 6)
new_point_2 = np.array([9, 6])
try:
```

```

#Broadcasting is not possible here.
#X is (4, 3) and the new point is (2,n).
#The last dimensions (3 and 2) are not equal and neither is 1
which violates the broadcasting rules.
result_2 = X + new_point_2
print("\nCase 2: Adding a point with two dimensions (9, 6)")
print("Broadcasting is possible. Result:")
print(result_2)
except ValueError as e:
    print("\nCase 2: Adding a point with two dimensions (9, 6)")
    print(f"Broadcasting is not possible. Error: {e}")

#Case 3: A datapoint with one dimension (2)
new_point_3 = np.array([2])
try:
    #Broadcasting is possible here.
    #The new point, (1,n), is compatible with X,(4, 3), because the
    last dimension is 1.
    #The new point is stretched across all elements of X.
    result_3 = X + new_point_3
    print("\nCase 3: Adding a point with one dimension (2)")
    print("Broadcasting is possible. Result:")
    print(result_3)
except ValueError as e:
    print("\nCase 3: Adding a point with one dimension (2)")
    print(f"Broadcasting is not possible. Error: {e}")

```

Array X:

```

[[6 3 7]
 [4 6 9]
 [2 6 7]
 [4 3 7]]

```

Euclidean Distance Matrix:

```

[[0.    4.123 5.    2.   ]
 [4.123 0.    2.828 3.606]
 [5.    2.828 0.    3.606]
 [2.    3.606 3.606 0.   ]]

```

Discussion on Broadcasting:

Case 1: Adding a point with coordinates (3, 1, 2)

Broadcasting is possible. Result:

```

[[ 9  4  9]
 [ 7  7 11]
 [ 5  7  9]
 [ 7  4  9]]

```

Case 2: Adding a point with two dimensions (9, 6)

Broadcasting is not possible. Error: operands could not be broadcast together with shapes (4,3) (2,)

Case 3: Adding a point with one dimension (2)

Broadcasting is possible. Result:

```
[[ 8  5  9]
 [ 6  8 11]
 [ 4  8  9]
 [ 6  5  9]]
```